



Project Impact Report

UNICC AI Safety Lab: Impact Report



Student Name: Shoa Afroz

Program: MS in Management and Analytics

School: New York University, School of Professional Studies

Course: Applied Technical Project

Instructor: Dr. Steven W. Goss

Project Sponsor: United Nations International Computing Centre (UNICC)

Project Guide / Sponsor Representative: Dr. Andrés Fortino, Clinical Associate Professor, NYU SPS

Project Focus: AI Safety Evaluation System, Testing, UI/UX, and Deployment Readiness

Submission Date: May 2026

Impact Summary

The UNICC AI Safety Lab project created a working proof of concept for structured AI safety evaluation in a governance-sensitive environment. Its impact is demonstrated through a functioning Council of Experts prototype, a completed VeriMedia test run, expert-level risk outputs, a final consensus decision, audit logging, and a supporting Streamlit comparison layer.

The project's value is not only that a prototype was built. Its main contribution is that it shows how AI safety review can move from an informal or manual process into a repeatable, explainable, and evidence-based workflow. This directly supports UNICC's need for responsible AI testing, pre-deployment review, and governance-ready decision support.

Problem Addressed

AI systems are increasingly being tested for use in high-stakes environments, but evaluation processes often remain fragmented. A single reviewer, a one-time prompt, or a black-box model output is not enough for organizations that need traceability, consistency, and accountability.

For UNICC, the challenge is not only whether an AI system performs well. The larger question is whether it can be reviewed safely before deployment across dimensions such as security risk, bias, harmfulness, privacy, transparency, and governance compliance. This project addressed that gap by creating a structured evaluation workflow that turns AI safety concerns into a practical review process.

What Was Delivered

The project delivered a Council of Experts AI Safety Lab prototype. The system allows a reviewer to submit an AI system using either a GitHub repository URL or a text description. The submitted system is then evaluated by three expert modules:

1. Risk & Technical Standards Expert
2. AI Safety & Red-Team Expert
3. Governance & Regulatory Expert

Each expert produces a verdict, confidence level, reasoning, and cited sources. The system then moves into a deliberation phase where experts respond to one another's findings. Finally, the system produces a Council consensus decision that summarizes the final verdict, rationale, points of agreement, and points of dissent.

A supporting Streamlit dashboard was also developed to show how results could be presented in a more stakeholder-facing format through score-based outputs, downloadable JSON reports, and run comparison features.

Evidence of Impact

Impact Area	Evidence from Project	Why It Matters
Functional completion	The system accepted the VeriMedia GitHub repository and completed the evaluation workflow	Shows the project moved beyond concept into a working prototype
Multi-expert evaluation	Three expert modules evaluated the same AI system from different perspectives	Reduces dependence on a single black-box judgment
Consistent decision output	All three experts returned FLAG verdicts for VeriMedia	Shows alignment across independent evaluation lenses
Final consensus	The Council produced a final FLAG decision with 90 percent combined confidence	Demonstrates structured synthesis and decision support
Auditability	The audit log preserved input, expert outputs, deliberation, final decision, timestamp, and latency	Supports traceability and governance review
Stakeholder usability	Streamlit added score views, downloadable reports, and comparison outputs	Makes results easier for non-technical stakeholders to interpret

Demonstrated Value to UNICC

The project creates value for UNICC by acting as a potential pre-deployment safety gate. Before an AI system is moved forward, reviewers can use the workflow to identify risks, understand why those risks were flagged, and preserve evidence for future review.

This is especially relevant for UNICC because AI systems in institutional environments must be evaluated not only for technical performance but also for governance readiness. The prototype connects directly to this need by incorporating evaluation lenses informed by NIST AI RMF, OWASP LLM Top 10, MITRE ATLAS, ISO/IEC 42001, EU AI Act, UN, and UNESCO perspectives.

The VeriMedia test case shows this value clearly. The system did not simply approve or reject the repository. It flagged specific concerns related to privacy, bias, transparency, third-party model dependency, security, and governance risk. This makes the output useful for decision-makers because it identifies what requires further review before deployment.

Effectiveness of the Output

The project was effective because it met the core success expectation: it produced a structured, interpretable, and auditable AI safety evaluation workflow.

1. First, the workflow was completed end to end. The system accepted an input, fetched repository information, generated expert assessments, supported deliberation, and produced a final decision.
2. Second, the results were interpretable. The system did not provide only a score or verdict. It explained the reasoning behind each expert's assessment and showed how the final decision was reached.
3. Third, the system supported auditability. The full audit log made the evaluation process reviewable after execution, which is critical for governance-sensitive environments.
4. Fourth, the Streamlit comparison feature strengthened the project by showing how repeated runs could be reviewed for consistency and reproducibility. This is important because AI evaluation systems must be tested not only for output quality but also for stability across runs.

Practical Impact

The practical impact of this project is that it gives UNICC a usable starting point for operationalizing AI safety evaluation. Instead of relying on scattered documents, manual review, or unstructured AI judgments, the prototype offers a more disciplined workflow.

The system can help stakeholders:

- Identify safety and governance risks earlier
- Compare expert reasoning before making deployment decisions
- Preserve evidence for accountability
- Support structured conversations between technical and governance teams
- Translate AI safety frameworks into practical review steps

This makes the project relevant beyond the classroom. It reflects the type of applied solution organizations need as they adopt AI systems under increasing regulatory, ethical, and operational pressure.

Impact Boundary

The project should be viewed as a strong proof of concept rather than a production-ready enterprise platform. Its impact is demonstrated through prototype functionality, test evidence, and structured outputs, but future work is still needed before real deployment.

The system should be tested across more AI repositories, risk types, and organizational use cases. Prompt consistency, retrieval reliability, model stability, and run-to-run reproducibility should also be strengthened. These limitations are important because they define the realistic boundary of the current deliverable.

Even with these limitations, the project successfully proves that the Council of Experts model can make AI safety evaluation more transparent, repeatable, and useful for governance review.

Final Impact Statement

The UNICC AI Safety Lab project demonstrates measurable applied value through a working prototype, completed VeriMedia evaluation, aligned expert verdicts, a final Council decision with 90 percent combined confidence, audit logging, and a supporting comparison dashboard.

Its main impact is that it turns AI safety evaluation into a structured decision-support process. For UNICC, this provides a practical foundation for reviewing AI systems before deployment, identifying risks early, and supporting responsible AI governance with evidence rather than assumption.